

COP 4338 Class 17


Last Class:

Ch 8: Process Control: (check slides, example, and book)

Today's Class:

Continue Process Control and Threads!

Note: Second homework , worth 15% Due March 9th




□


□

□

Chapter 11. Threads



COP-4520



Taken from: Introduction to Parallel Computing

By: Dr. Liu

Modified by:
Leonardo Bobadilla

A good resource!


- I **strongly** encourage you to follow this tutorial:

<https://computing.llnl.gov/tutorials/pthreads/>


What Are threads?

- A thread:
 - Is an independent flow of control.
 - Operates within a process with other threads.

Monothreaded process



Multithreaded process



Threads Versus Processes

Process

memory state

File state

Thread

Processor

state

program counter

stack pointer

Registers

Thread

Processor

state

program counter

stack pointer

Registers

- Threads vs. Processes:
 - Threads use and exist within the process resources
 - A thread maintains its own stack and registers, scheduling properties, set of pending and blocked signals.
 - By default, threads' data are shared, processes' data are private.
- Secondary Threads vs. Initial Threads
 - An initial thread is created automatically when a process is created.
 - Secondary threads are peers.

Why Threads?

- To realize potential program performance gains:
 - On a uniprocessor, multi-threaded processes hide latency for memory access, I/O, and communication.
 - On a multiprocessor system, a process with multiple threads provides potential parallelism.
- Benefits of multithreaded programming (rather than using multiple processes):
 - Compared to the cost of creating and managing a process, a thread can be created and managed with much less operating system overhead.
 - All threads within a process share the same address space. Inter-thread communication is more efficient and than inter-process communication.

Threads Versus Processes

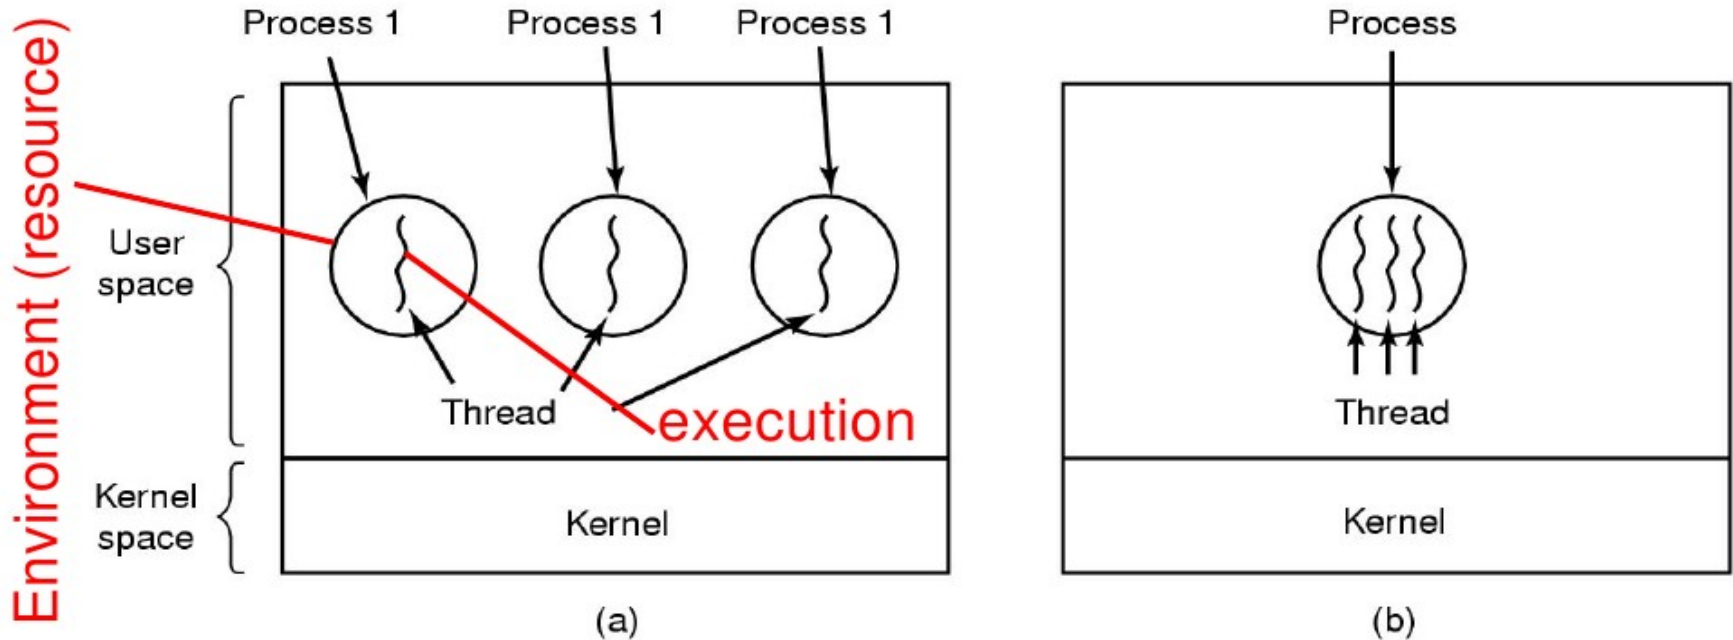
Process

fork is expensive (time & memory)

Thread

- Lightweight process
- Shared data space
- Does not require lots of memory or startup time

Threads Versus Processes



- a) Three processes each with one thread
- b) One process with three threads

What are POSIX threads (pthreads)?

- Historically, hardware vendors have implemented their own proprietary versions of threads. Not portable.
pthread is a standardized thread programming interface specified by the IEEE POSIX (portable operating systems interface) in 1995.
- pthreads are defined as a set of C language programming types and procedure calls, implemented with a header file (**pthread.h**) and a library (**libpthread.a**).

Outline

- Thread basics: creation and termination.
- Synchronization primitives:
 - Mutual exclusion.
 - Conditional variables.
- Thread cancellation.
- Composite synchronization constructs:
 - Read-write locks.
 - Barriers.

Thread Creation

- Creating a thread:

```
int pthread_create(  
    pthread_t *tid,                /* thread id */  
    const pthread_attr_t *attr,    /* thread attributes */  
    void *(*start_routine) (void *), /* pointer to function */  
    void *arg                      /* argument for function */  
);
```

- Returns the ID of the new thread via the *tid* argument.
- The *attr* parameter is used to set thread attributes, *NULL* as default.
- The *start_routine* is the C routine that the thread will start executing once it is created.
- A single argument may be passed to *start_routine* via *arg*. It must be passed by reference as a pointer cast of type void. You can use a data structure if you want to pass multiple arguments.

threadid.c threadexample1.c 14-threads.c simplechild1.c simplechild2.c
simplechild3.c

Thread Creation Arguments

Study `hello_arg1.c`

Study `hello_arg2.c`

Thread Termination

- There are many ways to terminate a thread:
 - The thread returns from its starting routine;
 - The thread calls **pthread_exit ()**;
 - The thread is canceled by another thread via **pthread_cancel()**;
 - The thread receives a signal that terminates it;
 - The entire process is terminated.

Thread Termination and Join

void pthread_exit (void *status) ;

- This function is used by a thread to terminate. The return value is passed as a pointer.

int pthread_join (pthread_t tid, void status);**

- The pthread_join() subroutine blocks the calling thread until the specified thread *tid* terminates.
- The function returns 0 on success, and negative on failure.
- The returned value is a pointer returned by reference. If you do not care about the return value, you can pass NULL for the second argument.

Hello World!

```
#include <pthread.h>
#include <stdio.h>
```

```
void *PrintHello(void * id) {
    printf("Thread%d: Hello World!\n", (int)id);
    pthread_exit(NULL); /* not necessary */
}
```

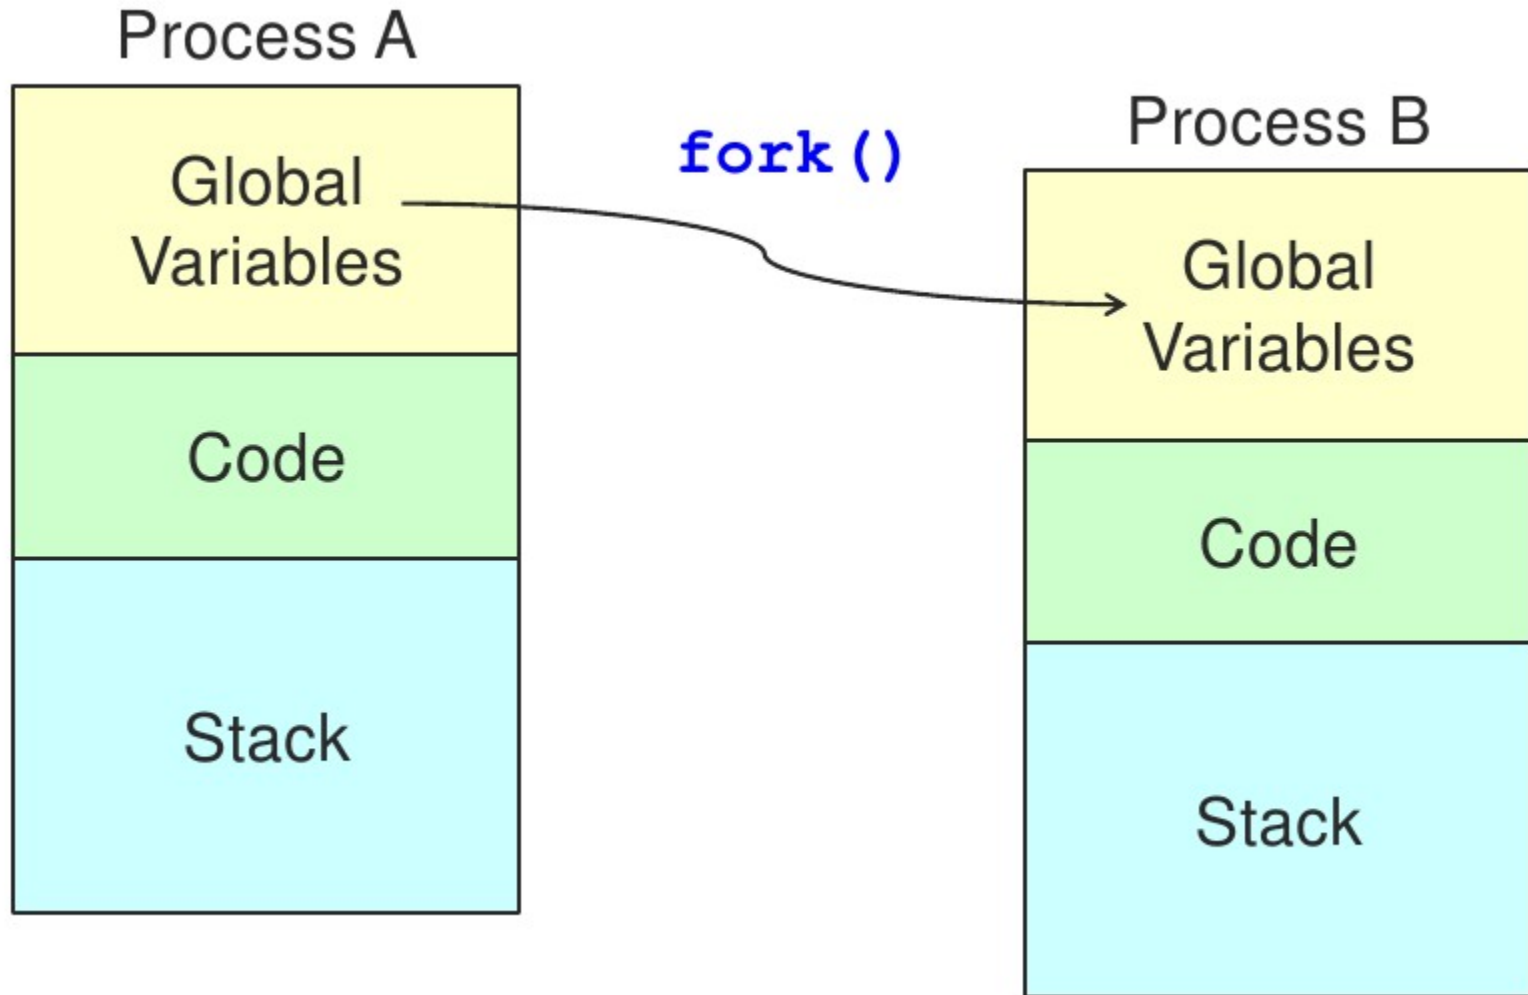
```
int main (int argc, char *argv[]) {
    pthread_t thread0, thread1;
    pthread_create(&thread0, NULL, PrintHello, (void *) 0);
    pthread_create(&thread1, NULL, PrintHello, (void *) 1);
    pthread_join(thread0, NULL);
    pthread_join(thread1, NULL);
    return 0;
}
```

See helloworldthreads.c pthreadjoinex.c

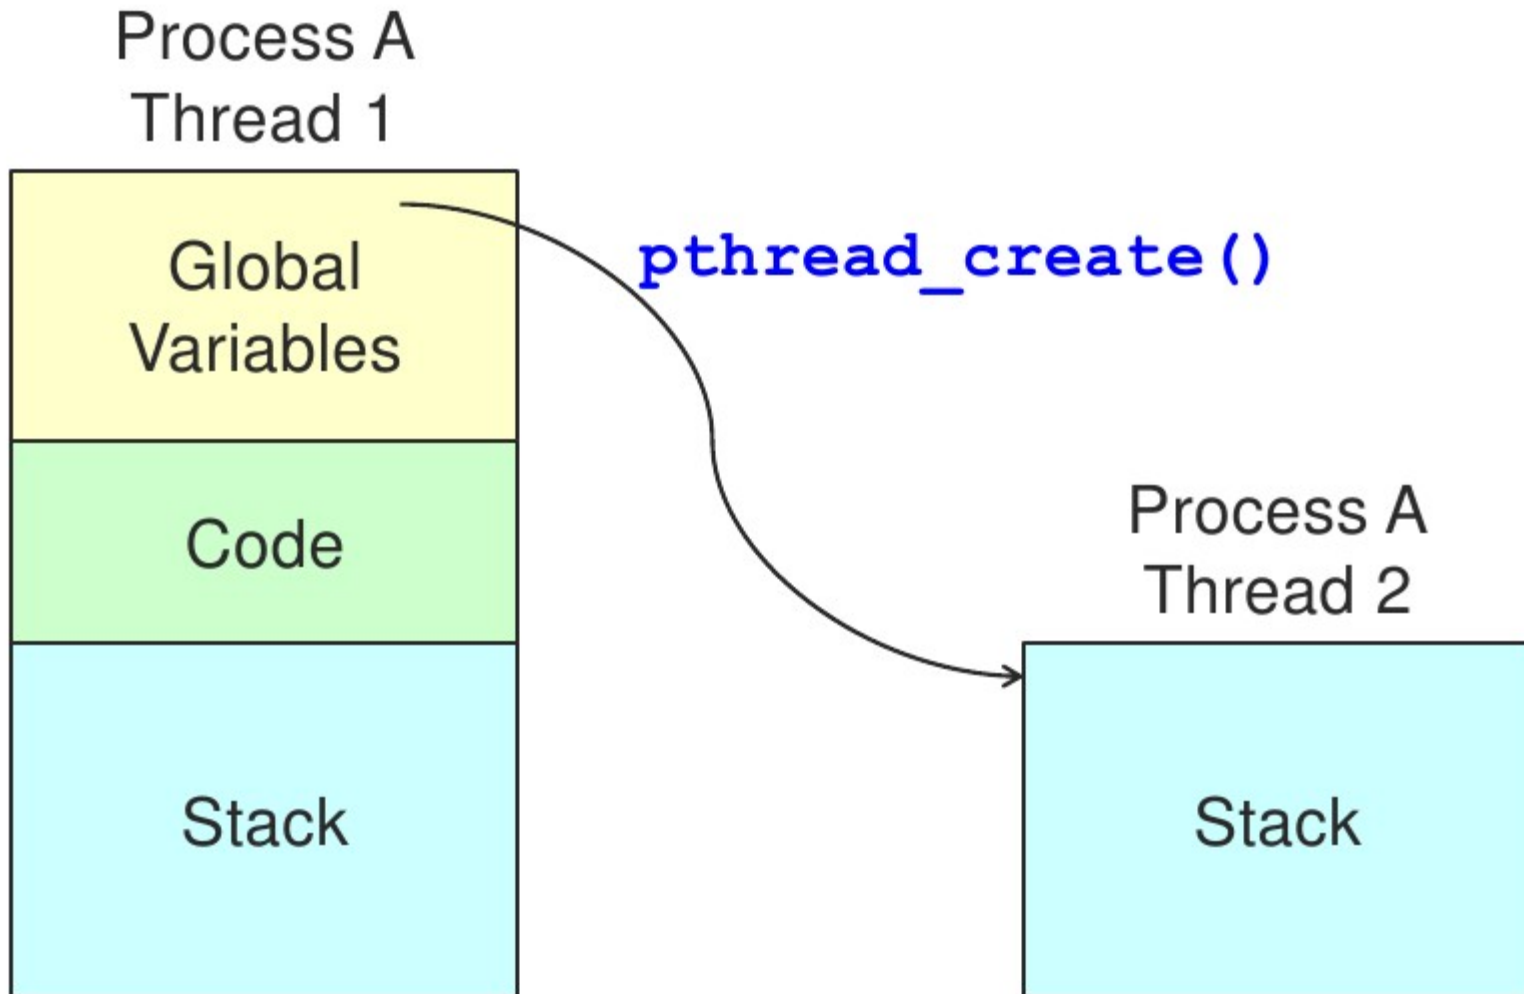
Thread vs Process creation

- **fork()** clones the process
 - Two separate processes with independent destinies
 - Independent memory space for each process
- **pthread_create()**
 - Start from a function
 - Share memory

Thread vs Process creation



Thread vs Process creation



Important Points

A thread is the lightest unit of work that can be scheduled to run on the processor

When creating a thread you

- Indicate which function the thread should execute

When a new thread is created

- It runs concurrently with the creating thread
- It shares common data space

Why threads over processes?

Creating a new process can be expensive

Time

- A call into the operating system is needed
- Context-switching involves the operating system

Memory

- The entire process must be replicated

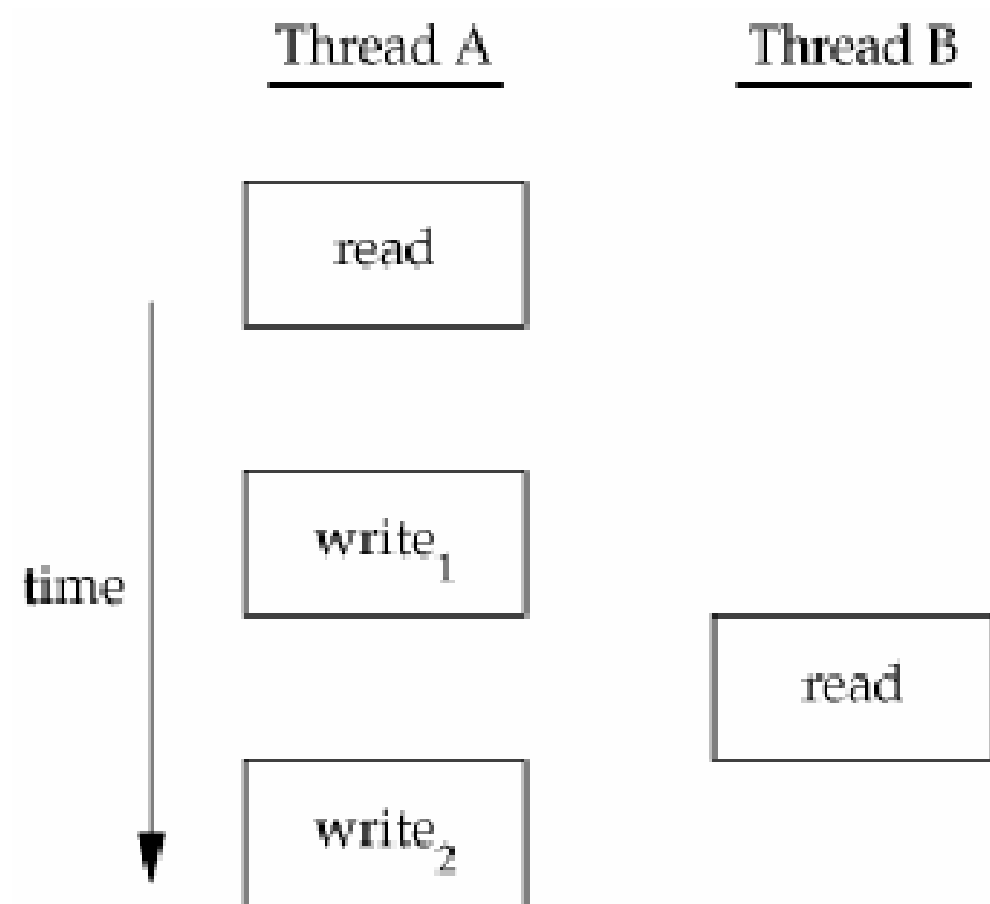
The cost of inter-process communication and synchronization of shared data

- May involve calls into the operation system kernel
- Threads can be created without replicating an entire process


Comparisson

Property	Processes created with fork	Threads of a process	Ordinary function calls
variables	Get copies of all variables	Share global variables	Share global variables
IDs	Get new process IDs	Share the same process ID but have unique thread ID	Share the same process ID (and thread ID)
Data/control	Must communicate explicitly, e.g., use pipes or small integer return value	May communicate with return value or carefully shared variables	May communicate with return value or shared variables
Parallelism (one CPU)	Concurrent	Concurrent	Sequential
Parallelism (multiple CPUs)	May be executed simultaneously	Kernel threads may be executed simultaneously	Sequential

Thread Synchronization



Thread Synchronization



Mutual Exclusion

- For thread synchronization and protecting shared data when multiple writes occur.
- A **mutex** variable acts like a "lock", protecting access to a shared data resource.
- Only one thread can obtain (or own) the lock at any given time. Thus, even if several threads try to lock a **mutex**, only one thread will be successful. No other thread can obtain the lock until the owning thread unlocks that mutex.
 - Threads must "take turns" accessing protected data.

Creating / Destroying Mutexes

```
pthread_mutex_t mutex;  
int pthread_mutex_init(pthread_mutex_t* mutex,  
    pthread_mutexattr_t* attr);  
int pthread_mutex_destroy (pthread_mutex_t* mutex);
```

- Mutex variables must be declared and initialized before they can be used.
- Usually, one specifies *attr*, the mutex object attributes, to be NULL to accept defaults.
- It's a good habit to destroy the mutex after use.

Locking / Unlocking Mutexes

```
int pthread_mutex_lock(pthread_mutex_t* mutex);  
int pthread_mutex_trylock(pthread_mutex_t* mutex);  
int pthread_mutex_unlock(pthread_mutex_t* mutex);
```

- The `pthread_mutex_lock()` routine is used by a thread to acquire a lock on the specified *mutex* variable. The thread **blocks** if the mutex is already locked by another thread.
- `pthread_mutex_trylock()` will attempt to lock a mutex. However, if the mutex is already locked, the routine will **return immediately** with a "busy" error code.
- `pthread_mutex_unlock()` will unlock a mutex; must be called by the owning thread.

A Typical Scenario

```
pthread_mutex_t mutex;
```

```
int main() {  
    pthread_mutex_init(&mutex, NULL);  
    pthread_create(&t0, NULL, p, NULL);  
    pthread_create(&t1, NULL, q, NULL);  
    ...  
    pthread_join(t0, NULL);  
    pthread_join(t1, NULL);  
    pthread_mutex_destroy(&mutex);  
    return 0;  
}
```

syncexample.c mutex1.c bank.c

```
void p(void* arg) {  
    ...  
    pthread_mutex_lock(&mutex);  
    /* access shared variable */  
    pthread_mutex_unlock(&mutex);  
    ...  
}  
void q(void* arg) {  
    ...  
    pthread_mutex_lock(&mutex);  
    /* access shared variable */  
    pthread_mutex_unlock(&mutex);  
    ...  
}
```

Deadlocks

- Deadlocks occur when a thread try to relock a mutex it already owns:

```
f() {  
    pthread_mutex_lock ( & mutex);  
    ...  
    g();  
    ...  
    pthread_mutex_unlock();  
}  
  
g() {  
    pthread_mutex_lock ( & mutex);  
    ...  
    pthread_mutex_unlock();  
}
```

- Deadlocks may also occur when mutexes are locked in reverse order:

<pre>/* Thread A */ pthread_mutex_lock(&mutex1); time 0 pthread_mutex_lock(&mutex2); time 2</pre>	<pre>/* Thread B */ pthread_mutex_lock(&mutex2); time1 pthread_mutex_lock(&mutex1); time3</pre>
---	---

Notes: To avoid this, successive mutexes should be locked in the same order.

Basic Rules on Using Mutexes

- Avoid deadlocks:
 - If it must use 2 locks simultaneously, be careful of the locking order.
- Improve performance:
 - Move all unnecessary code outside the critical section; the code inside a critical section is doomed to be sequential.
 - Accesses to shared data should be put together if they can be covered by a single lock/unlock.

Condition Variables

- Condition variables allow threads to synchronize based upon the actual data value;
- Without condition variables, the thread has to check the condition continuously (i.e., polling). This can be very resource consuming;
- With condition variables, the thread sleeps and waits on the condition variable;
- A condition variable is always used in conjunction with a mutex.

Condition Variables

standard "waiter" idiom:

```
pthread_mutex_lock( &mutex );  
while ( !condition )  
    pthread_cond_wait( &cond, &mutex );  
do_something();  
pthread_mutex_unlock( &mutex );
```

This function atomically releases mutex and causes the “waiter” to block on the condition variable “cond”. Upon successfully return, the mutex has been locked and owned by the “waiter” again.

standard "signaler" idiom:

```
pthread_mutex_lock( &mutex );  
do_something();  
// make condition TRUE  
pthread_cond_signal( &cond );  
pthread_mutex_unlock( &mutex );
```

If more than one thread is blocked on “cond”, the scheduling policy determines the order in which threads are unblocked. `pthread_cond_broadcast ()` unblocks all threads currently blocked on “cond”.

Notes: `pthread_cond_wait()` should be called while mutex is locked;
`pthread_cond_signal()` should be called after mutex is locked;
It is a logical error to call signal before calling wait.

Condition Variables

Check condvar.c
cond1.c

Classical Producer-Consumer Problem

- The producer produces data that are consumed by the consumer.



- The producer will only be able to produce a data while the queue has available space.
- The consumer will only be able to consume a data while the queue has data in it.

```

pthread_mutex_t mutex;
pthread_cond_t cond_empty, cond_full;
    queue state */
int data_available = 0;

main() {
    pthread_mutex_init(&mutex, NULL);
    pthread_cond_init(&cond_empty, NULL);
    pthread_cond_init(&cond_full, NULL);
}

```

```

void producer(void* arg) {
    for(;;) {
        produce_data();
        pthread_mutex_lock(&mutex);
        while(data_available)
            data */
        pthread_cond_wait(&cond_empty, &mutex);
        insert_data_into_queue();
        space */
        data_available = 1;
        retrieve */
        pthread_cond_signal(&cond_full);
        pthread_mutex_unlock(&mutex);
    }
}

```

```

void consumer(void* arg) {
    for(;;) {
        pthread_mutex_lock(&mutex);
        while(!data_available)
            queue */
        pthread_cond_wait(&cond_full, &mutex);
        retrieve_data_from_queue();
        must be one */
        data_available = 0;
        pthread_cond_signal(&cond_empty);
        pthread_mutex_unlock(&mutex);
        consume_data();
    }
}

```

```

/* declare mutex */
/* declare two conditional variables for different

/* declare and initialize global variables */

```

```

/* initialize the mutex */
/* initialize the condition variables */

```

```

/* produce the data first, outside the critical section */
/* will block until the lock is obtained */
/* IMPORTANT: check whether there's space to insert

/* if not, wait until there is space */
/* insert the data into the queue, knowing there's

/* set the condition that the data is now available for

/* time to wake up the consumer */
/* get out of the critical section */

```

```

/* try to enter the critical section */
/* IMPORTANT: check whether there's data in the

/* if not, wait until there's data */
/* retrieve the data from the queue, knowing there

/* set the condition that the queue is now empty */
/* time to wake up the producer to produce more */
/* get out of the critical section */
/* consume the data, OUTSIDE of the critical section

```

Example

run pc_1.c

Thread Cancellation

- A thread may cancel itself or other threads.
- It is not guaranteed that the specified thread will receive or act on the cancellation.
- If cancellation is acted upon, it will terminate the thread as if **pthread_exit()** is called.

int pthread_cancel(pthread_t thread);

Composite Synchronization Constructs

- POSIX threads API only provides a basic set of synchronization primitives: mutex, condition variables, etc.
- We can use these synchronization primitives to build higher level constructs.
- Read-Write Locks & Barriers (check that in the book)

Composite Synchronization Constructs

- Suppose A , B , and C are $N \times N$ matrices, we are suppose to calculate $C = A \times B$.
- Suppose we have p^2 number of threads (i.e., the number of threads is a square number) and we assume N is divisible by p .

Composite Synchronization Constructs

- We can multiply A and B block by block. For explanation, let's assume $p=4$. That is, we have 16 threads. Matrix A can thus be divided into 4×4 submatrices, as follows:

A00 A01 A02 A03

A10 A11 A12 A13

A20 A21 A22 A23

A30 A31 A32 A33

Composite Synchronization Constructs

- You can do the same for matrix B and matrix C.
- Now, to calculate matrix C from A and B. We can have each thread to be in charge of a submatrix of C. Let's assume it's the i th row and j th column (where $0 \leq i \leq 3$, and $0 \leq j \leq 3$ in this example): $C_{ij} = A_{i0} \times B_{0j} + A_{i1} \times B_{1j} + A_{i2} \times B_{2j} + A_{i3} \times B_{3j}$

References

The previous and following slides are mostly taken from:

<http://www.cs.stevens.edu/~jschauma/810D/>

