

# **Parallel Computing**

# Why parallel computing?

# Why parallel computing?

[https://computing.llnl.gov/tutorials/parallel\\_comp/](https://computing.llnl.gov/tutorials/parallel_comp/)

# Why parallel computing?

Please install Message Passing Interface MPI  
in your computer!

<http://jetcracker.wordpress.com/2012/03/01/how-to>

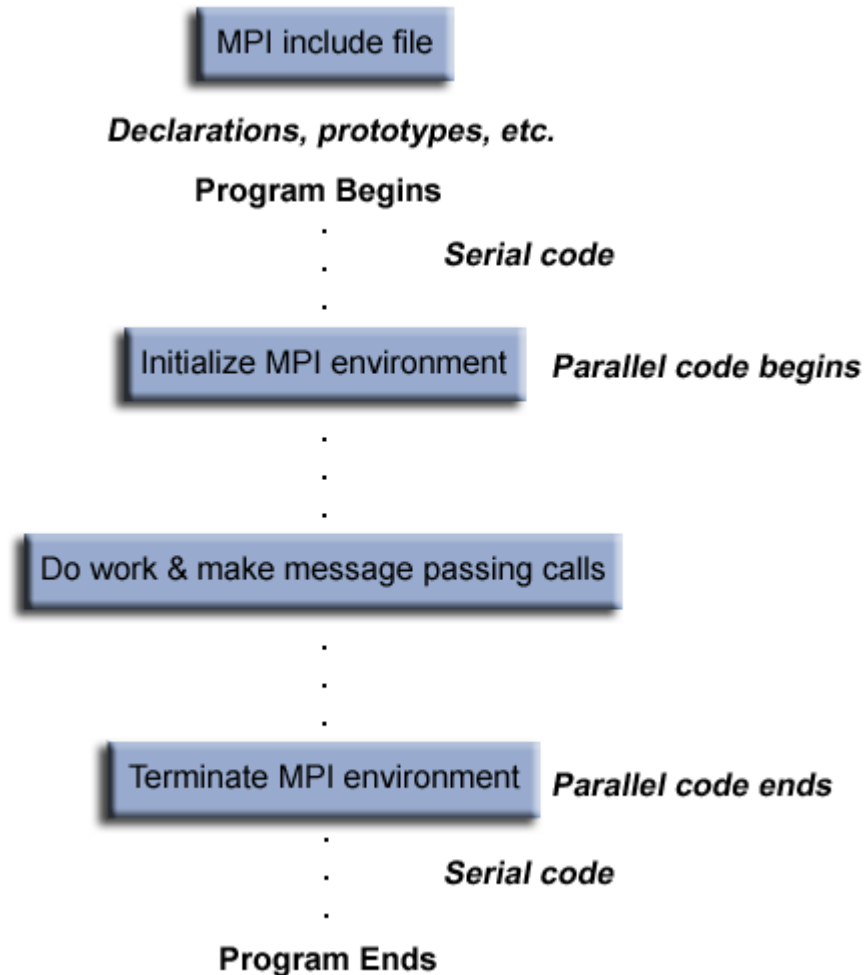
<http://en.wikipedia.org/wiki/MPICH>

Read more about MPI:

<https://computing.llnl.gov/tutorials/mpi/>

[http://en.wikipedia.org/wiki/Message\\_Passing\\_Inte](http://en.wikipedia.org/wiki/Message_Passing_Interface)

# Structure of the program



# MPI

mpi-hello.c

mpi-example1.c

mpi-sumschem.c

mpi-sum.c

# What is MPI?

- *A message-passing library specification*
  - extended message-passing model
  - not a language or compiler specification
  - not a specific implementation or product
- For parallel computers, clusters, and heterogeneous networks
- Full-featured
- Designed to provide access to advanced parallel hardware for
  - end users
  - library writers
  - tool developers

# Why Use MPI?

- MPI provides a powerful, efficient, and *portable* way to express parallel programs
- MPI was explicitly designed to enable libraries...
- ... which may eliminate the need for many users to learn (much of) MPI




# Why Use MPI?

- Standardization - MPI is the only message passing library which can be considered a standard. It is supported on virtually all HPC platforms. Practically, it has replaced all previous message passing libraries.
- Portability - There is little or no need to modify your source code when you port your application to a different platform that supports (and is compliant with) the MPI standard.
- Performance Opportunities - Vendor implementations should be able to exploit native hardware features to optimize performance.
- Functionality - There are over 440 routines defined in MPI-3, which includes the majority of those in MPI-2 and MPI-1.
- Availability - A variety of implementations are available, both vendor and public domain.

# **Programming with MPI**

# Message-passing Model



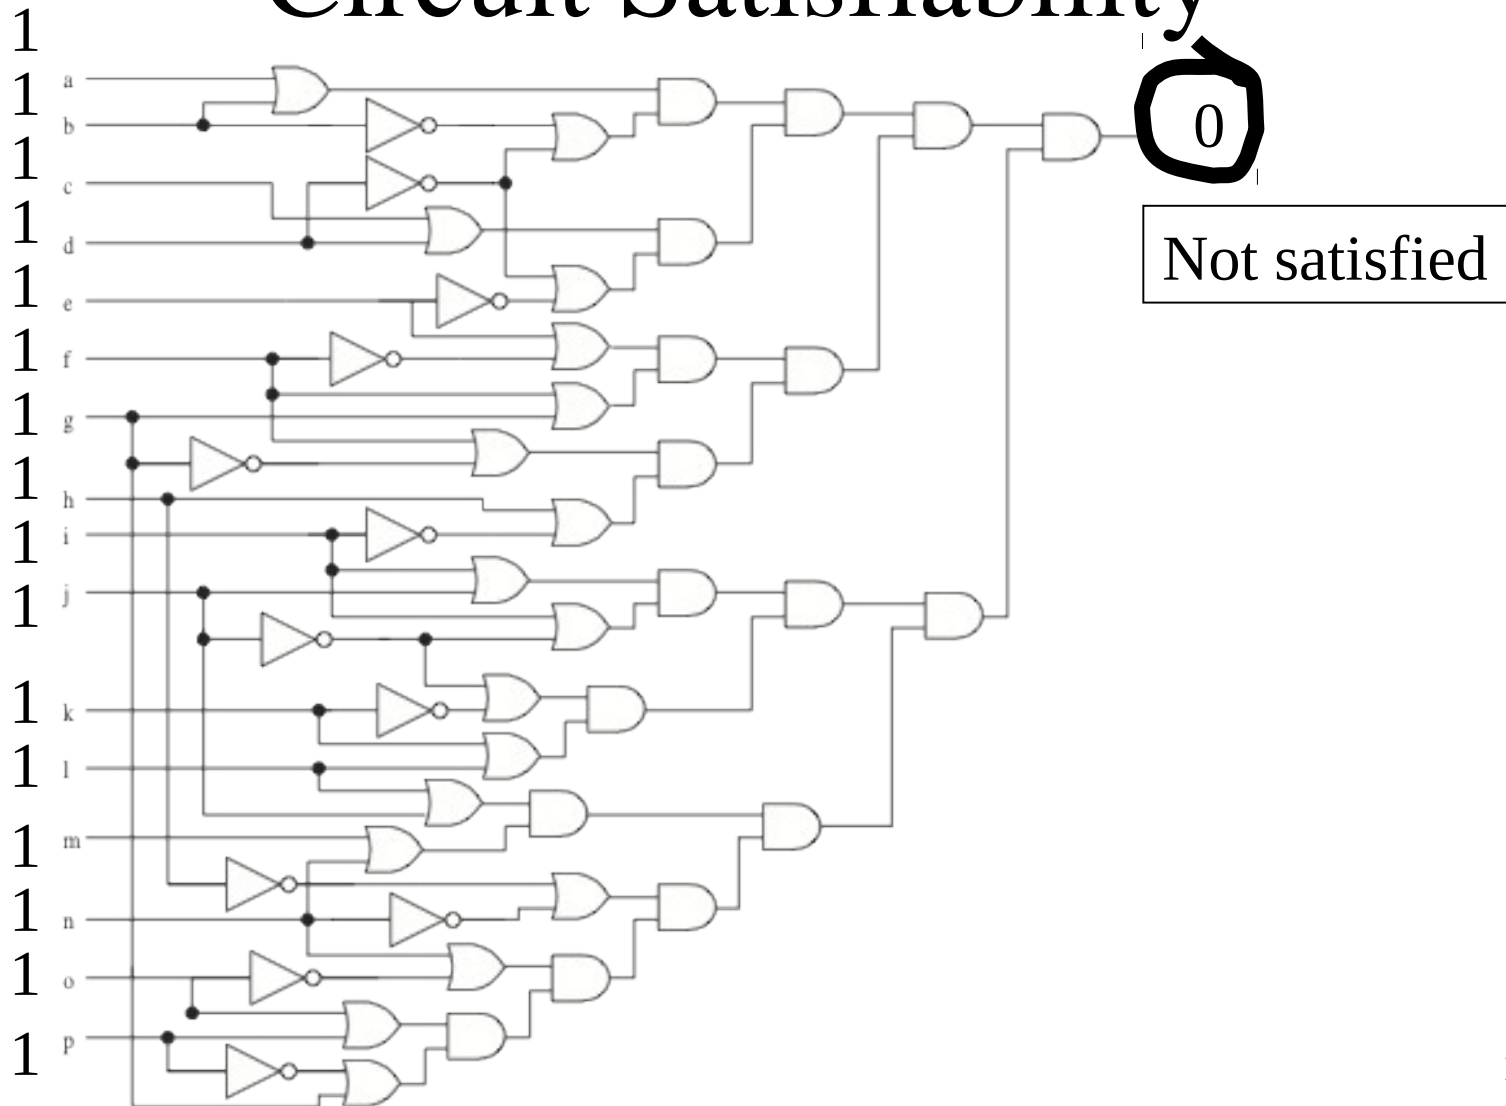
# Processes

- Number is specified at start-up time.
  - Typically, fixed throughout the execution.
- All execute same program (SPMD).
  - Each distinguished by a unique ID number.
- Processes explicitly pass messages to communicate and to synchronize with each other.

# Advantages of Message-passing Model

- Gives programmer ability to manage the memory hierarchy.
  - What's local, what's not?
- Portability to many architectures: can run both on shared-memory and distributed-memory platforms.

# Circuit Satisfiability



```

/* Return 1 if 'i'th bit of 'n' is 1; 0 otherwise */
#define EXTRACT_BIT(n,i) ((n&(1<<i))>0)

void check_circuit (int id, int z) {
    int v[16];          /* Each element is a bit of z */
    int i;

    for (i = 0; i < 16; i++) v[i] = EXTRACT_BIT(z,i);

    if ((v[0] || v[1]) && (!v[1] || !v[3]) && (v[2] || v[3])
        && (!v[3] || !v[4]) && (v[4] || !v[5])
        && (v[5] || !v[6]) && (v[5] || v[6])
        && (v[6] || !v[15]) && (v[7] || !v[8])
        && (!v[7] || !v[13]) && (v[8] || v[9])
        && (v[8] || !v[9]) && (!v[9] || !v[10])
        && (v[9] || v[11]) && (v[10] || v[11])
        && (v[12] || v[13]) && (v[13] || !v[14])
        && (v[14] || v[15])) {
        printf ("%d) %d%d%d%d%d%d%d%d%d%d%d%d%d\n", id,
            v[0],v[1],v[2],v[3],v[4],v[5],v[6],v[7],v[8],v[9],
            v[10],v[11],v[12],v[13],v[14],v[15]);
        fflush (stdout);
    }
}

```

# Solution Method

- Circuit satisfiability is NP-complete.
- No known algorithms to solve in polynomial time.
- We seek all solutions.
- We find through exhaustive search.
- 16 inputs  $\Rightarrow$  65,536 combinations to test.



# A parallel version: Summary of Program Design

- Program will consider all 65,536 combinations of 16 boolean inputs.
- Combinations allocated in cyclic fashion to processes.
- Each process examines each of its combinations.
- If it finds a satisfiable combination, it will print it.

# Include Files

**#include <mpi.h>**

- MPI header file.

**#include <stdio.h>**

- Standard I/O header file.

# Local Variables

```
int main (int argc, char *argv[]) {  
    int i;  
    int id; /* Process rank */  
    int p; /* Number of processes */
```

- Include **argc** and **argv**: they are needed to initialize MPI
- One copy of every variable for each process running this program

# Initialize MPI

**MPI\_Init (&argc, &argv);**

- First MPI function called by each process.
- Not necessarily first executable statement.
- Allows system to do any necessary setup.

# Shutting Down MPI

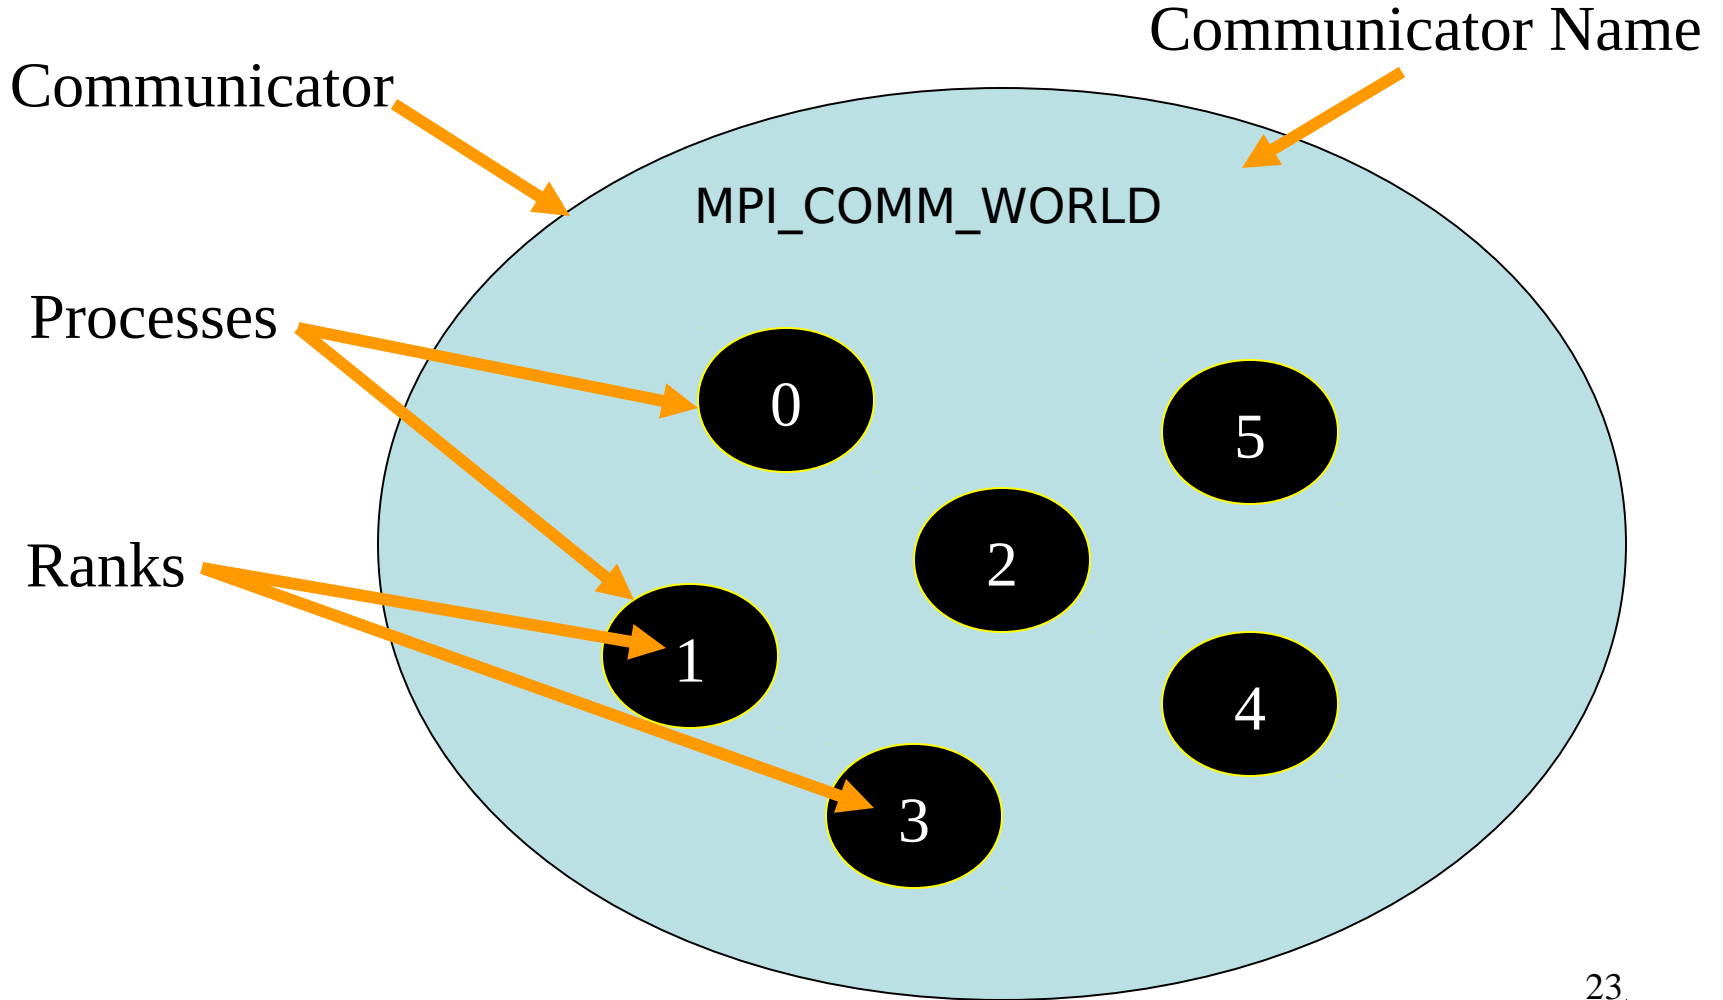
## **MPI\_Finalize();**

- Call after all other MPI library calls
- Allows system to free up MPI resources

# Communicators

- Communicator: opaque object that provides message-passing environment for processes.
- `MPI_COMM_WORLD`
  - Default communicator.
  - Includes all processes that participate the run.
- It's possible to create new communicators by user.
  - Always a subset of processes defined in the default communicator.

# Communicator



# Determine Number of Processes

**MPI\_Comm\_size (MPI\_COMM\_WORLD, &p);**

- First argument is communicator,
- Number of processes returned through second argument.




# Determine Process Rank

**MPI\_Comm\_rank (MPI\_COMM\_WORLD, &id);**

- First argument is communicator.
- Process rank (in range 0, 1, ...,  $p-1$ ) returned through second argument.

# Replication of Automatic Variables



# Cyclic Allocation of Work

```
for (i = id; i < 65536; i += p)  
  check_circuit (id, i);
```

- Parallelism is outside function **check\_circuit**
- It can be an ordinary, sequential function

```
#include <mpi.h>
#include <stdio.h>
```

```
int main (int argc, char *argv[]) {
    int i;
    int id;
    int p;
```

```
    MPI_Init (&argc, &argv);
    MPI_Comm_rank (MPI_COMM_WORLD, &id);
    MPI_Comm_size (MPI_COMM_WORLD, &p);
```

```
    for (i = id; i < 65536; i += p)
        check_circuit (id, i);
```

```
    printf ("Process %d is done\n", id);
    fflush (stdout);
    MPI_Finalize();
    return 0;
```

```
}
```

**Put fflush() after every printf()**

# Compiling MPI Programs

```
% mpicc -O -o foo foo.c
```

- **mpicc**: script to compile and link MPI programs.
- Flags: same meaning as C compiler.
  - **-O** — optimization level.
  - **-o <file>** — where to put executable.

# Running MPI Programs

**% mpirun -np <p> <exec> <arg1> ...**

- **-np <p>** — number of processes.
- **<exec>** — executable.
- **<arg1> ...** — command-line arguments.

# Execution on 1 CPU

```
% mpirun -np 1 sat
0) 10101111110011001
0) 01101111110011001
0) 11101111110011001
0) 1010111111011001
0) 0110111111011001
0) 1110111111011001
0) 10101111110111001
0) 01101111110111001
0) 11101111110111001
Process 0 is done
```

# Execution on 2 CPUs

```
% mpirun -np 2 sat
0) 0110111110011001
0) 011011111011001
0) 0110111110111001
1) 1010111110011001
1) 1110111110011001
1) 101011111011001
1) 111011111011001
1) 1010111110111001
1) 1110111110111001
Process 0 is done
Process 1 is done
```



# Execution on 3 CPUs

```
% mpirun -np 3 sat
0) 0110111110011001
0) 1110111111011001
2) 1010111110011001
1) 1110111110011001
1) 1010111111011001
1) 0110111110111001
0) 1010111110111001
2) 0110111111011001
2) 1110111110111001
Process 1 is done
Process 2 is done
Process 0 is done
```

# Deciphering Output

- Output order only partially reflects order of output events inside parallel computer.
- If process A prints two messages, first message will appear before second.
- If process A calls **printf** before process B, there is no guarantee process A's message will appear before process B's message.

# Enhancing the Program

- We want to find total number of solutions.
- Incorporate sum-reduction into program.
- Reduction is a **collective communication operation**.

# Modifications

- Modify function **check\_circuit**
  - Return 1 if circuit satisfiable with input combination.
  - Return 0 otherwise.
- Each process keeps local count of satisfiable circuits it has found.
- Perform reduction after **for** loop.

# New Declarations and Code

```
int count; /* Local sum */  
int global_count; /* Global sum */  
  
count = 0;  
for (i = id; i < 65536; i += p)  
    count += check_circuit (id, i);
```

# Prototype of **MPI\_Reduce()**

```
int MPI_Reduce (  
    void *operand, /* addr of 1st reduction element */  
    void *result,  /* addr of 1st reduction result */  
    int count,     /* reductions to perform */  
    MPI_Datatype type, /* type of elements */  
    MPI_Op operator, /* reduction operator */  
    int root,      /* process getting result(s) */  
    MPI_Comm comm  /* communicator */  
);
```

# MPI\_Datatype Options

- **MPI\_CHAR**
- **MPI\_DOUBLE**
- **MPI\_FLOAT**
- **MPI\_INT**
- **MPI\_LONG**
- **MPI\_LONG\_DOUBLE**
- **MPI\_SHORT**
- **MPI\_UNSIGNED\_CHAR**
- **MPI\_UNSIGNED**
- **MPI\_UNSIGNED\_LONG**
- **MPI\_UNSIGNED\_SHORT**

# MPI\_Op Options

- **MPI\_BAND**
- **MPI BOR**
- **MPI\_BXOR**
- **MPI LAND**
- **MPI\_LOR**
- **MPI\_LXOR**
- **MPI\_MAX**
- **MPI\_MAXLOC**
- **MPI\_MIN**
- **MPI\_MINLOC**
- **MPI\_PROD**
- **MPI\_SUM**



# Our Call to **MPI\_Reduce()**

```
MPI_Reduce (&count,  
            &global_count,  
            1,  
            MPI_INT,  
            MPI_SUM,  
            0,  
            MPI_COMM_WORLD);
```

Only process 0  
will get the result

```
if (!id) printf ("There are %d different solutions\n",  
                global_count);
```

# Execution of Second Program

```
% mpirun -np 3 seq2
```

```
0) 01101111110011001
```

```
0) 11101111111011001
```

```
1) 11101111110011001
```

```
1) 10101111111011001
```

```
2) 10101111110011001
```

```
2) 01101111111011001
```

```
2) 11101111110111001
```

```
1) 01101111110111001
```

```
0) 10101111110111001
```

```
Process 1 is done
```

```
Process 2 is done
```

```
Process 0 is done
```

```
There are 9 different solutions
```

# Benchmarking the Program

- **MPI\_Barrier** — barrier synchronization.
- **MPI\_Wtick** — timer resolution.
- **MPI\_Wtime** — current time.

# Benchmarking Code

```
double elapsed_time;  
...  
MPI_Init (&argc, &argv);  
MPI_Barrier (MPI_COMM_WORLD);  
elapsed_time = - MPI_Wtime();  
...  
MPI_Reduce (...);  
elapsed_time += MPI_Wtime();
```

# MPI Functions So Far

- **Basic Functions:**

- **MPI\_Init**
- **MPI\_Comm\_rank**
- **MPI\_Comm\_size**
- **MPI\_Finalize**

- **More:**

- **MPI\_Reduce**
- **MPI\_Barrier**
- **MPI\_Wtime**
- **MPI\_Wtick**

# Two More Functions to Conquer the World

```
int MPI_Send(
    void* buf,                /* addr of 1st element to send */
    int count,                /* number of elements */
    MPI_Datatype datatype,    /* type of elements */
    int dest,                 /* destination process id */
    int tag,                  /* message tag */
    MPI_Comm comm             /* communicator */
);

int MPI_Recv(
    void* buf,                /* addr of 1st element to receive */
    int count,                /* number of elements */
    MPI_Datatype datatype,    /* type of elements */
    int source,               /* source process id */
    int tag,                  /* message tag */
    MPI_Comm comm,            /* communicator */
    MPI_Status* status        /* info about message received */
);
```

# More on **MPI\_Recv**

- **MPI\_Recv** is a blocking function.
- Use wildcards:
  - **MPI\_ANY\_SOURCE.**
  - **MPI\_ANY\_TAG.**
- You can check the status of the received message:
  - MPI\_Status.**MPI\_SOURCE**
  - MPI\_Status.**MPI\_TAG**